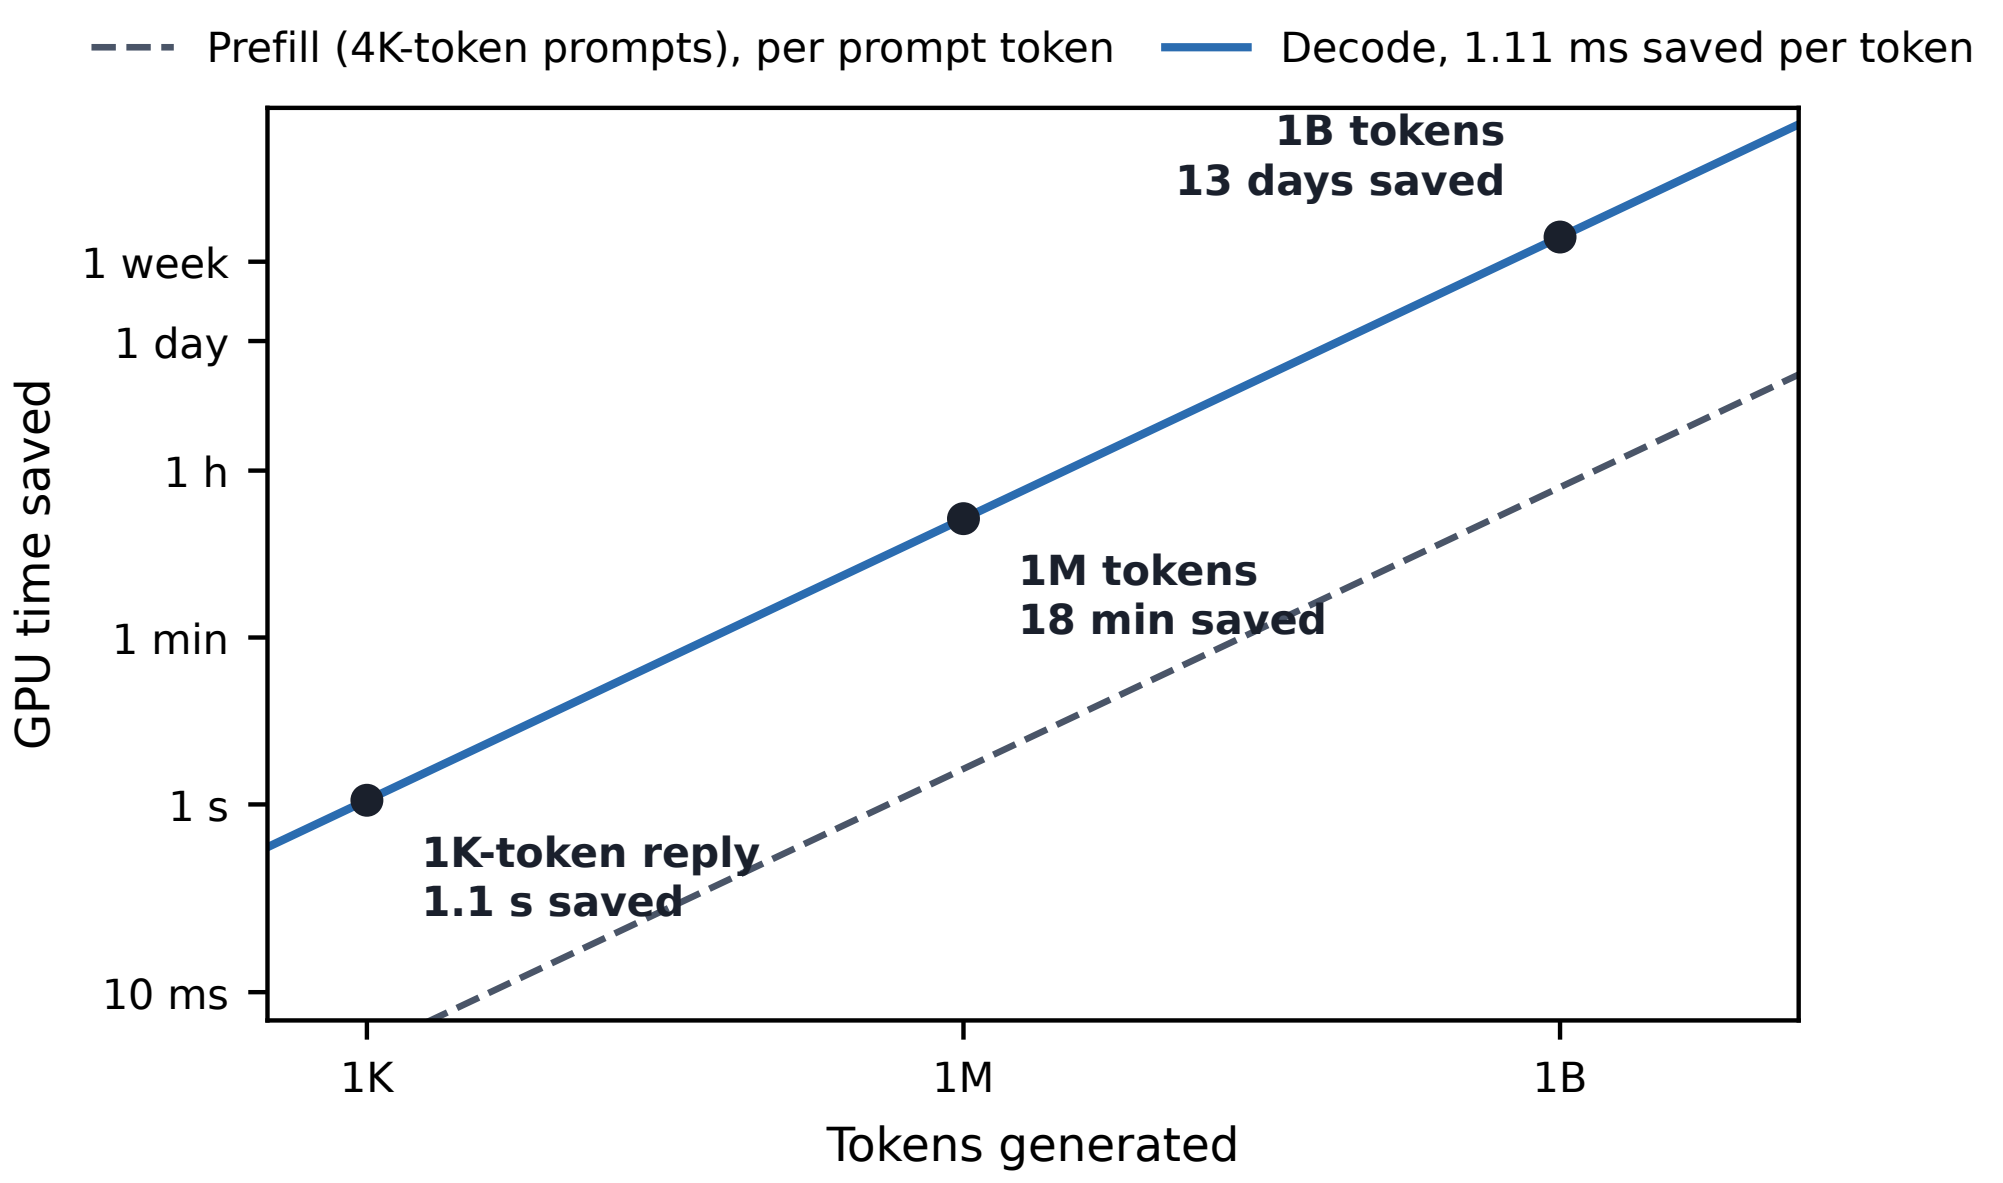
Prefill (4K-token prompts), per prompt token
Decode, 1.11 ms saved per token
1B tokens
13 days saved
1M tokens
18 min saved
1K-token reply
1.1 s saved
1 week
1 day
1 h
1 min
1 s
10 ms
GPU time saved
1K
1M
1B
Tokens generated